


Prompting barely changes outcomes for a reasoning model (GSM8K, n = 200 per condition)

model global:deepseek-v4.1-flash, temperature 0, 600 generations

(a) Accuracy by condition



● exact match ○ after numeric correction


Intervals overlap completely: no condition separates from another.

Grading: exact match under-counts by 1.0 points — 6 responses were numerically equal but textually different (12 vs 12.00), 2 failed to parse.

Item difficulty and the grading rule dominate; the prompt does not. 188 of 200 items were right in all three conditions and 8 in none.

Data: GSM8K test split (Cobbe et al., 2021). aidoc/experiment-cot-gsm8k.json. Every number recomputed from the stored generations.

(b) Per-item outcome, Direct vs CoT



(c) Answer vs hidden reasoning

